

User Story 01 — Demand Intake Agent

Scope of this document:

- Defines the MVP user story for a Foundry-hosted Demand Intake Agent.
- Covers two customer-facing channels: Microsoft Teams and embedded WebChat in the WebApp.
- Focuses on direct customer demand submission, adaptive guided intake, validation, duplicate detection, customer confirmation, record creation, status visibility, and audit logging.
- Excludes downstream scoring, governance triage, CPQ creation, routing, and delivery allocation except where required as integration handoff points.

1. User Story Summary

Epic	Demand Management — Customer-facing Intake
Feature	Foundry-hosted Demand Intake Agent available through Teams and WebChat
Primary user	Customer user, such as transformation director, product owner, business sponsor, or authorised demand submitter
Secondary users	NTT DATA governance lead, NTT DATA platform administrator, support analyst
Release	MVP
Priority	Must have
Channels	Microsoft Teams bot/channel app and embedded WebChat in the AI Studio Delivery Command Center WebApp
Hosting pattern	Azure AI Foundry-hosted conversational agent with secure API integration to the platform data layer
User Story Overview	I want to submit a new AI Studio demand item through a conversational agent in either Microsoft Teams or the WebApp WebChat, so that I can describe my request in plain language, be guided to provide the right information, check whether similar demand already exists, confirm the final captured record, and immediately track the submitted item without relying on email, forms, or an NTT DATA account team to manually raise it for me.

2. Business Context

The Demand Intake Agent is the customer-facing front door for the AI Studio Delivery Command Center. The MVP requires customers to raise demand directly through the platform rather than relying on account teams to relay requests. The agent must guide the customer through intake, adapt the question flow to the demand type, validate completeness, perform semantic duplicate detection, and create a governed demand item that becomes visible in the customer's tracker with status **Submitted**.

This story supports the wider Demand Management module by creating a clean, structured demand item that can then be scored, prioritised, triaged, and routed by NTT DATA users. The customer must only see safe customer-facing fields and must never see internal scoring, commercial detail, capacity data, or governance rationale.

3. Goals and Outcomes

Goal	Outcome
Remove manual intake dependency	Customers can raise demand themselves through Teams or WebChat without sending emails or asking an account manager to create the request.

Capture complete structured demand records	The agent collects mandatory and conditional information needed for downstream triage and scoring.
Improve submission quality	The agent asks clarifying questions, validates missing fields, and keeps incomplete submissions in draft rather than creating low-quality demand records.
Reduce duplicate demand	The agent checks semantically similar demand items, including other business units within the same customer organisation, before final submission.
Provide immediate customer visibility	Once confirmed, the demand item appears in the customer tracker with status Submitted and a customer-safe summary.
Support governance and auditability	Every agent action, question, answer, duplicate match, customer confirmation, and record creation event is logged for audit and governance review.

4. Personas

Persona	Needs	Permissions
Customer submitter	Raise a new demand item, answer guided questions, review duplicate suggestions, confirm submission, track status.	Can create and view own organisation's demand items based on customer access rules. Cannot view internal scores, CPQ, rates, margins, or capacity.
Customer approver / sponsor	May submit strategic demand or validate business value and sponsorship details.	Can create demand and view organisation-level demand if provisioned with that role.
NTT DATA Solution Lead	Review created demand items, inspect agent-captured context, view duplicate flags, continue triage.	Can view all demand fields, duplicate rationale, internal notes, scoring inputs, and audit trail.
Platform Administrator	Configure question flow, mandatory fields, demand types, thresholds, channel settings, fallback messages.	Can configure metadata, prompts, routing endpoints, and operational settings.


5. Scope

MVP Scope — Demand Intake AgentFoundry-hosted agent · Microsoft Teams · WebChat · Guided intake · Duplicate detection · Draft handling · Submission confirmation · Demand record creation

- Azure AI Foundry-hosted Demand Intake Agent.
- Teams channel for customer demand intake.
- Embedded WebChat channel within the WebApp.
- Adaptive conversational intake by demand type.
- Mandatory and conditional field validation.
- Semantic duplicate detection across customer organisation demand backlog.
- Customer review and confirmation before record creation.
- Draft state for incomplete submissions.
- Agent action logging and traceability.

6. Main User Journey

The intake journey must work consistently across Teams and Webchat. Channel-specific user experience may differ, but the business outcome, captured data, validation rules, duplicate logic, and created demand record must be the same.



7. Detailed Functional Requirements

FR-001 — Agent availability across two channels

The same Demand Intake Agent capability must be available in Microsoft Teams and embedded Webchat.

The customer user must be able to start a new intake conversation from either Microsoft Teams or the AI Studio Delivery Command Center WebApp.

The Webchat should be embedded within the customer portal and should recognise the logged-in user's organisation context. The Teams channel should support customer users who have been provisioned to access the agent through approved tenant or external access setup.

Expected behaviour: A user who starts in Teams and a user who starts in Webchat are guided through equivalent intake logic. The created demand record must have the same data structure regardless of channel.

FR-002 — User authentication and organisation context

The agent must only allow authorised customer users to submit demand for their own organisation.

Before a demand record is created, the platform must resolve the user's identity, organisation, permitted business units, and customer role. If the user's identity or organisation cannot be resolved, the agent must not create a demand item. It should provide a controlled message instructing the user to contact their platform administrator or NTT DATA support contact.

Security behaviour: The agent must not ask the user to manually type their organisation as the sole source of truth if authenticated organisation context is available. User-provided organisation text can be captured for clarification, but record ownership must come from the platform identity context.

FR-003 — Natural language demand capture

The user can begin by describing their need in plain language.

The agent must accept a free-text description such as "We want an AI tool to review supplier contracts" or "We need additional data engineering capacity for our chatbot build". The agent should extract an initial summary, potential demand type, business area, likely outcomes, and missing information.

In practice: The user writes a rough request. The agent responds with a short acknowledgement, proposes a demand type, and asks the next best question rather than forcing the user into a long static form.

FR-004 — Demand type classification

The agent adapts the question flow based on the type of demand.

Demand type	Examples	Additional question focus
AI use case / new solution	New AI assistant, document intelligence tool, predictive model, workflow transformation idea.	Business problem, target users, expected value, process impacted, data availability, timeline, sponsor.
Enhancement to existing solution	Improve existing chatbot, add features to current analytics product, expand model coverage.	Existing solution name, current limitation, requested change, affected users, urgency, production impact.
Capacity request	Need additional developers, architects, data engineers, testers, or delivery support.	Required roles, skill areas, start date, duration, expected commitment, location or shore preference.
Exploratory / unclear request	Customer has a broad idea but does not know if it is a use case, enhancement, or capacity ask.	Clarification questions; agent may keep classification as Unknown until enough information is captured.

FR-005 — Mandatory field capture

The agent must collect the minimum data required to create a demand item.

Field	Required?	Notes
Demand title	Yes	Agent may generate suggested title from conversation and allow customer edit.
Demand type	Yes	AI use case, enhancement, capacity request, or other configured type.
Description	Yes	Plain language description of request, generated or confirmed from conversation.
Business area	Yes	Customer business unit, department, function, or process area.
Business problem / objective	Yes	What the customer is trying to solve or achieve.
Expected value or benefit	Yes, where known	Can be financial, efficiency, risk reduction, customer experience, compliance, or qualitative benefit.
Proposed timeline	Yes	Target start, desired completion, urgency, or key business date.
Submitter details	Yes	Derived from authenticated user profile; not manually trusted as source of identity.
Customer organisation	Yes	Derived from access context.
Attachments or supporting links	Optional	May be captured as references if enabled by channel and platform policy.

Consent to submit	Yes	User must confirm final summary before creation.
-------------------	-----	--

FR-006 — Conditional field capture

The agent asks extra questions when the demand type requires them.

Condition	Additional fields
Capacity request	Role type, skill requirements, number of people or estimated effort, required start date, duration, location preference, urgency driver.
Enhancement to existing solution	Existing solution name, current owner if known, existing issue or limitation, affected users, release or deadline impact.
New AI use case	Target process, data sources if known, expected users, decision or workflow impacted, risk/compliance sensitivity, desired outcome.
Potential regulated or sensitive data use	Data sensitivity indicator, personal data indicator, compliance or security notes for NTT DATA review.
Unclear or low-confidence classification	Agent asks clarification questions and may flag the item for governance review as ambiguous.

FR-007 — Completeness validation and draft handling

Incomplete requests must not become unactionable records.

The agent must validate whether mandatory fields are present before creating a demand item. If required information is missing, the agent should continue asking focused questions. If the user cannot complete the submission, the agent should save the conversation as a draft demand intake where supported, mark the draft as incomplete, and allow the user to resume later.

Expected behaviour: The platform must not create a submitted demand item with no business area, no description, or no timeline. Incomplete items remain draft and are visible only as draft/intake-in-progress where customer visibility rules allow.

FR-008 — Semantic duplicate detection

The agent checks for similar existing demand items before final submission.

Before creating a new demand item, the agent must search existing demand records using semantic similarity, not keyword matching alone. The search must include demand items from the same customer organisation and, where access rules permit, other business units within that organisation. The agent must identify exact duplicates, likely duplicates, and related-but-distinct demand items.

Agent behaviour: The agent must never auto-merge demand records. It should show a customer-safe summary of potential matches, ask whether the new request is the same or distinct, and record the user's choice. Similar but distinct items should be cross-referenced on creation.

Duplicate result	Customer-facing behaviour	Record behaviour
High-confidence duplicate	Agent shows customer-safe summary of existing item and asks whether the user wants to follow existing item or continue as distinct.	If user confirms same initiative, no new demand item is created unless business rules require a linked submission record.
Possible duplicate	Agent explains that a similar request exists and asks the user to confirm same vs distinct.	If distinct, new demand item is created with semantic duplicate reference and flagged for governance review.
Related demand	Agent shows that related activity exists but does not block submission.	New demand item is created with a related-demand cross-reference.
No match	Agent proceeds to final summary and confirmation.	New demand item is created normally.

FR-009 — Customer review and confirmation

The user must confirm the structured demand summary before submission.

Before creating the demand item, the agent must present a concise summary containing title, demand type, business area, objective, expected value, proposed timeline, known risks or constraints, and duplicate/related-demand outcome. The user must be able to edit incorrect fields or confirm submission.

Confirmation options: Confirm and submit; edit a field; add more detail; cancel; save as draft.

FR-010 — Demand record creation

A confirmed submission creates a governed Demand Item record.

Once the customer confirms the summary, the platform must create a Demand Item record in the data layer. The record must include the captured business fields, customer organisation, submitter ID, channel of submission, duplicate references, generated summary, raw conversation reference, status, and audit metadata.

Initial status:The internal status should be created as Submitted or equivalent configured backend status. The customer-facing status label must display as **Submitted**.

FR-011 — Customer acknowledgement and tracker visibility

The customer immediately sees the submitted item.

After record creation, the agent must confirm successful submission and provide the customer with the demand reference ID, title, submitted date /time, current customer-facing status, and where to track it. In Webchat, a link to the demand tracker should be shown. In Teams, a card or message should provide the same safe summary and link if enabled.

Customer sees:Reference ID, title, status Submitted, high-level summary, and next-step expectation such as "NTT DATA will review this request." Customer must not see score, scoring breakdown, governance notes, CPQ fields, rates, margin, or capacity data.

FR-012 — Agent action logging

Every material agent action must be auditable.

The agent must write entries to the Agent Action Log for conversation start, identity resolution, demand type classification, questions asked, answers received, completeness checks, duplicate search execution, duplicate results shown, customer confirmation, record creation, and any error or fallback event.

Audit behaviour:Logs must support traceability without exposing unnecessary customer-sensitive data to users who are not authorised to view it.

8. Channel-specific Requirements

Area	Microsoft Teams	Webchat
Entry point	User opens the Demand Intake Agent in Teams, through an approved app, chat, or channel experience.	User opens embedded chat from the customer portal or demand submission page.
Authentication	Identity resolved through Teams/Microsoft identity context and platform provisioning.	Identity resolved through WebApp session and platform RBAC.
Conversation UX	Conversational messages, adaptive cards where appropriate, quick reply buttons for options.	Conversational panel with structured cards, field summary, edit controls, and tracker link.
Draft resume	Agent should support resume where user identity and conversation state can be linked.	Draft should be accessible from customer portal draft/intake-in-progress area where implemented.
Submission confirmation	Teams confirmation card/message with demand reference and status.	Webchat confirmation with direct link to submitted demand tracker detail page.
Attachments	Attachment support depends on approved Teams and platform file handling policy.	Attachment or link capture depends on WebApp upload controls and security policy.
Error handling	Controlled Teams message with retry or support path.	Inline Webchat error message with retry and support path.

9. Data Requirements

Data object	Required fields for this story	Notes
Demand Item	Demand ID, title, type, description, business area, objective, expected value, proposed timeline, status, customer-visible status label, submitter ID, customer organisation ID, source channel, duplicate references, created timestamp.	Core record created by this story.
Demand Intake Draft	Draft ID, user ID, organisation ID, partial answers, missing fields, last updated timestamp, channel, expiry policy.	Required if the user cannot complete all mandatory information.
Duplicate Match Result	Search ID, candidate demand ID, similarity score band, match type, customer-safe summary shown, user decision, timestamp.	Supports audit and governance review.
Agent Action Log	Log ID, agent name, action type, input reference, output reference, timestamp, channel, user ID, confidence score where relevant, human confirmation status.	Mandatory audit trail.
Conversation Transcript Reference	Conversation ID, channel, storage reference, retention category, related demand ID or draft ID.	Should store reference or governed transcript according to data policy.

10. Acceptance Criteria

#	Scenario	Given	When	Then
AC-001	Start intake in Teams	An authorised customer user has access to the Teams Demand Intake Agent	The user starts a new demand conversation	The agent greets the user, resolves their identity and organisation context, and asks for the demand description.
AC-002	Start intake in Webchat	An authorised customer user is logged into the WebApp	The user opens the Demand Intake WebChat	The agent starts the same intake flow using the WebApp session context.
AC-003	Unauthorised user blocked	A user cannot be matched to an authorised customer organisation	The user attempts to submit demand	The agent prevents record creation and displays a controlled access message.
AC-004	Demand type inferred	The user describes a new AI solution request	The agent processes the description	The agent proposes the demand type and asks relevant follow-up questions.
AC-005	Capacity request flow	The user asks for additional delivery capacity	The agent classifies the request as capacity-related	The agent asks for role, skill, start date, duration, and location or shore preference.
AC-006	Mandatory fields missing	The user has not provided business area or proposed timeline	The user tries to submit	The agent explains what is missing and asks focused follow-up questions before allowing submission.
AC-007	Save incomplete draft	The user cannot provide all mandatory information	The user chooses to stop or save progress	The agent saves an incomplete draft where supported and does not create a Submitted demand item.
AC-008	Semantic duplicate found	A similar demand item exists in the same customer organisation	The agent performs duplicate detection	The agent shows a customer-safe summary of the similar item and asks whether this is the same or distinct.
AC-009	User confirms distinct demand	A possible duplicate has been shown	The user confirms the new request is distinct	The platform creates a new demand item with a duplicate or related-demand reference for governance review.
AC-010	User confirms same demand	A high-confidence duplicate has been shown	The user confirms it is the same initiative	The platform does not create an unnecessary duplicate demand item and guides the user to the existing customer-visible item where permitted.
AC-011	Final review required	The agent has collected all required information	The user reaches submission step	The agent displays a structured summary and requires explicit confirmation before creation.
AC-012	Edit before submission	The final summary contains an incorrect field	The user chooses to edit	The agent updates the relevant field and redisplay the summary before confirmation.
AC-013	Demand item created	The user confirms the final summary	The platform saves the demand item	A Demand Item record is created with status Submitted and the correct customer-visible status label.
AC-014	Customer acknowledgement	A demand item has been created successfully	The agent responds to the user	The user receives a demand reference ID, title, status Submitted, and tracker guidance or link.
AC-015	Customer visibility restriction	A demand item has been submitted	The customer views the item	The customer can see only customer-safe fields and cannot see internal score, scoring rationale, CPQ, rate, margin, or capacity information.
AC-016	Agent action log written	The intake conversation has completed or failed	The platform records the process	Agent actions are logged with channel, user, action type, timestamp, and related demand or draft reference.
AC-017	Channel consistency	The same demand information is provided via Teams and WebChat in separate tests	Each user confirms submission	The resulting Demand Item records follow the same schema and lifecycle rules regardless of channel.
AC-018	Error handling	The demand creation API fails	The user confirms submission	The agent shows a controlled failure message, does not claim success, logs the error, and allows retry or support escalation.

11. Non-functional Requirements

Category	Requirement
Security	All access must respect customer and NTT DATA role-based access controls. Customer users must only create and view demand within their authorised organisation scope.
Privacy	Conversation transcripts, attachments, and extracted summaries must be handled according to approved data retention and privacy policies.
Reliability	The agent must not create duplicate records due to retry, refresh, or repeated confirmation events. Record creation must be idempotent where technically feasible.
Availability	The intake capability should remain available in at least one channel if the other channel is temporarily unavailable, subject to platform architecture.
Performance	Common agent responses should feel interactive. Duplicate search and record creation should return clear progress messages if they take longer than expected.

Explainability	The agent should explain why it is asking for key fields in plain language, especially value, timeline, and business area.
Auditability	All material actions must be traceable through Agent Action Log and linked to the demand item or draft.
Accessibility	Webchat must support keyboard navigation, readable contrast, and screen-reader-friendly labels where applicable.
Branding	The Webchat and confirmation components should follow the AI Studio / NTT DATA Command Center visual language used by the platform.

12. Error and Edge Case Handling

Case	Expected handling
User abandons conversation	Save as draft if enough context exists and draft capability is enabled; otherwise log abandoned session without creating a demand item.
User provides vague request	Agent asks clarifying questions and may classify as exploratory or unknown until enough detail is gathered.
User asks for unsupported action	Agent explains that it can help submit demand and redirects to supported intake path.
User tries to view internal scoring or commercial detail	Agent refuses to expose internal information and provides customer-safe status information only.
Duplicate search unavailable	Agent should not silently skip duplicate detection unless product rules permit. It should either retry, save draft, or create with a flag that duplicate check failed and needs governance review.
Record creation fails	Agent must not claim the item was submitted. It should show retry/support guidance and log the failure.
User changes channel mid-process	If conversation state is linked to the user and draft, the user should be able to resume. If not supported, agent should start a new intake and avoid exposing previous channel data unsafely.
Multiple users submit similar demand	Semantic duplicate logic should surface possible matches and capture whether the request is same, related, or distinct.
User enters sensitive or regulated data	Agent should capture a sensitivity flag and avoid unnecessary repetition of sensitive content in summaries.

13. Integration Requirements

Integration	Purpose	Requirement
Azure AI Foundry	Agent hosting, orchestration, prompt/tool execution, model governance path.	Demand Intake Agent must be deployed as a governed Foundry-hosted agent or equivalent approved architecture.
Microsoft Teams	Customer-facing conversational channel.	Agent must be accessible to approved customer users through Teams subject to tenant and access setup.
WebApp Webchat	Embedded customer portal channel.	Agent must be embedded in the WebApp and use authenticated WebApp context.
Platform API	Demand record creation and draft management.	Agent must create and update demand records through secure APIs, not by direct unauthorised database writes.
Data layer / Fabric	Demand items, drafts, duplicate search index, agent action logs.	Created records and audit entries must persist to the approved data model.
Semantic search / vector index	Duplicate and related-demand detection.	Agent must search existing demand records using semantic similarity and customer access boundaries.
Notification service	Submission acknowledgement and future status updates.	Agent should trigger customer-safe acknowledgement and support tracker visibility after creation.

14. Sample Conversation

Example — New AI use case

Customer: We want to build an AI tool for contract review.
Agent: I can help capture that as a new AI Studio demand item. Is this a new AI use case, an enhancement to an existing solution, or a request for additional delivery capacity?
Customer: New AI use case.
Agent: Great. Which business area owns this request?
Customer: Legal Operations.
Agent: What problem are you trying to solve?
Customer: Contract review takes too long and we want to identify risky clauses faster.
Agent: What benefit are you expecting?
Customer: Faster review times and fewer missed compliance issues.
Agent: Is there a target timeline or business deadline?
Customer: We would like a discovery to start in September.
Agent: I found a similar request from another business unit in your organisation: "Automated legal document analysis" submitted three weeks ago. Is your request the same initiative or a separate but related request?
Customer: Separate, but related.
Agent: Thanks. Here is the demand summary I will submit...
Customer: Confirm and submit.
Agent: Submitted. Your demand reference is DEM-000123. Current status: Submitted. You can track it in your demand tracker.

15. Definition of Done

#	Criterion
1	The Foundry-hosted Demand Intake Agent is available through Microsoft Teams and embedded WebChat.
2	Authorised customer users can start a demand intake conversation from either channel.
3	The agent resolves user identity and customer organisation context before creating a demand item.
4	The agent captures demand type, description, business area, objective, expected value or benefit, and proposed timeline.
5	The agent adapts its questions for AI use case, enhancement, capacity request, and unclear demand types.
6	Mandatory field validation prevents incomplete demand items from being submitted.
7	Incomplete intake can be saved as draft where supported and is not treated as submitted demand.
8	Semantic duplicate detection runs before final submission and checks relevant demand records within the customer organisation boundary.
9	The agent shows customer-safe duplicate or related-demand summaries and captures the user's decision.
10	The customer reviews and explicitly confirms the final structured summary before record creation.
11	A confirmed submission creates a Demand Item record with customer-visible status Submitted.
12	The customer receives a demand reference and can see the submitted item in their demand tracker.
13	Customer views do not expose internal score, scoring rationale, CPQ, rate card, margin, resource, or capacity data.
14	Agent actions are written to the Agent Action Log with sufficient traceability for governance review.
15	Failed duplicate search, failed record creation, access issues, and abandoned conversations are handled with controlled user messages and logs.
16	Both channels create demand records using the same platform API and data schema.
17	The created demand item can be picked up by the downstream Demand Scoring Agent or triage queue after submission.

16. Guided Questions Catalogue

The Demand Intake Agent should use the following question bank dynamically rather than asking every question in every conversation. The goal is to capture enough information to create a complete demand item and link it to measurable benefits, ROI, value evidence, and business outcomes from the point of submission.

Questioning principles

The agent should ask short, plain-language questions, adapt follow-ups based on the demand type, and avoid forcing the customer to provide financial precision when proxy data can support a value estimate.

Principle	Expected behaviour
Start broad, then qualify	Begin with the user's problem or idea, then ask structured follow-up questions only where needed.
Capture value early	Ask about benefit category, affected volume, time spent, cost, risk, revenue, and target KPI before submission.
Use proxy ROI questions	If the user cannot estimate financial value, ask about transaction volume, manual effort, affected users, error rate, SLA impact, or revenue/risk exposure.
Do not expose internal scoring	The agent may capture information used for prioritisation, but must not tell the customer how NTT DATA scoring will be calculated.
Make duplicate handling transparent	When similar demand is found, show a customer-safe summary and ask whether the new request is the same, related, or separate.
Confirm before creation	The agent must present a structured demand summary and obtain explicit confirmation before creating the submitted demand item.

A. Demand type and context

Used to classify the request, associate it with the right organisation and business area, and establish ownership.

#	Guided question	Purpose / data captured
1	What type of request are you raising?	Classifies demand type. Suggested options: new AI use case, enhancement to existing solution, automation opportunity, additional capacity request, advisory/discovery request, or other.
2	Which business area, function, or process does this demand relate to?	Links the demand to a business domain for triage, reporting, and future portfolio grouping.
3	Which customer organisation, business unit, region, or operating company should this demand be associated with?	Supports customer-scoped visibility, duplicate detection, and cross-business-unit related-demand checks.
4	Who is the business owner or sponsor for this demand?	Identifies accountable sponsor and enables follow-up.
5	Who should NTT DATA contact if clarification is needed?	Identifies day-to-day contact for the more-information-needed workflow.
6	Is this request related to an existing project, product, engagement, or previous demand item?	Helps link related records and avoid duplicate effort.
7	Are there any supporting links, documents, examples, screenshots, or current process notes you want to add?	Captures supporting evidence or attachments for d

B. Problem statement

Used to understand the current-state pain, affected users, frequency, and urgency.

#	Guided question	Purpose / data captured
8	What problem or opportunity are you trying to address?	Captures the core business need in plain language.
9	What is happening today that makes this difficult, slow, costly, risky, or inefficient?	Establishes current-state pain and reason for change.
10	Who is affected by this problem?	Identifies impacted users, teams, customers, operations, or external stakeholders.
11	How often does this problem occur?	Captures frequency, scale, and repeatability.
12	How many cases, transactions, documents, calls, tickets, decisions, or process instances are affected per month?	Creates a volume baseline for ROI and benefit estimation.
13	What happens if this demand is not addressed?	Captures risk of doing nothing and supports urgency assessment.
14	Is there a deadline, event, regulatory driver, customer commitment, or business milestone behind this request?	Captures time sensitivity and business urgency.

C. Desired outcome and benefits

Used to link the demand to measurable outcomes and business benefits.

#	Guided question	Purpose / data captured
15	What outcome are you expecting if this demand is delivered successfully?	Defines the target state and expected result.
16	Which benefits do you expect this demand to create?	Allows benefit tagging. Suggested options: cost reduction, time saving, revenue growth, risk reduction, quality improvement, customer experience improvement, employee productivity, compliance, innovation, scalability.
17	What would success look like in measurable terms?	Creates success criteria for later benefit tracking.
18	Which KPI, SLA, business metric, or operational measure should improve?	Links the demand to a measurable business metric.
19	What is the current baseline for that metric, if known?	Captures current performance before delivery.
20	What target improvement would you like to achieve?	Captures expected uplift or improvement target.
21	How soon do you expect the benefits to be realised after delivery?	Captures benefit timing and ROI horizon.
22	How will the benefit be measured after the solution is live?	Defines measurement approach for benefit realisation.

D. ROI and value hypothesis

Used to capture the financial, operational, strategic, risk, or customer-value case linked to the demand.

#	Guided question	Purpose / data captured
23	What is the estimated annual value of solving this problem?	Captures customer-stated annual value estimate where known.
24	Is the value mainly financial, operational, strategic, regulatory, or customer-experience related?	Classifies value type for prioritisation and reporting.
25	Can you estimate the current cost of the problem?	Supports ROI calculation and cost-of-inaction view.
26	How much time is currently spent on this activity per transaction, case, document, or cycle?	Enables productivity-saving estimate.
27	Which roles or teams spend time on this today?	Supports cost proxy calculations without asking for personal salary data.
28	Approximately how many people are involved in the current process?	Quantifies scale of manual effort and affected user base.
29	Could this demand help generate new revenue, accelerate sales, improve conversion, or protect existing revenue?	Captures revenue-side value.
30	Could this demand reduce penalties, leakage, rework, failed SLAs, compliance exposure, or operational risk?	Captures risk avoidance and value protection.
31	Are there external costs today, such as third-party review, manual processing, vendor support, or consultancy spend?	Identifies direct cost-reduction opportunities.
32	Is there a known budget, funding source, or business case already attached to this demand?	Helps assess commercial readiness and routing.
33	What confidence level do you have in the value estimate?	Captures estimate reliability. Suggested options: low, medium, high.
34	What evidence supports the benefit estimate?	Captures value evidence such as operational reports, volumes, SLA reports, cost data, or stakeholder estimates.

E. AI, automation, and data fit

Used to identify the likely solution pattern and early feasibility considerations.

#	Guided question	Purpose / data captured
35	What kind of AI or automation capability do you think may be needed?	Captures initial solution hypothesis. Suggested options: generative AI, document intelligence, chatbot, knowledge search, workflow automation, predictive analytics, agentic automation, or not sure.
36	What data, documents, systems, or knowledge sources would the solution need to use?	Identifies data dependencies.
37	Is the required data already available and accessible?	Captures early feasibility signal.
38	Who owns the data or system access needed for this demand?	Identifies data owner and future dependency.
39	Are there any sensitive, regulated, confidential, or personal data considerations?	Captures security, privacy, regulatory, and governance flags.
40	Does this request require changes to an existing business workflow or process?	Captures future workflow transformation indicator without making Layer 2 part of the MVP.
41	Is this a standalone solution, or does it need to integrate with existing systems?	Captures integration complexity.
42	Are there existing tools, pilots, spreadsheets, macros, manual workarounds, or prototypes related to this demand?	Identifies existing assets and avoids rebuilding known solutions.

F. Delivery expectations

Used to understand timeline, delivery level, dependencies, and constraints.

#	Guided question	Purpose / data captured
43	When would you ideally like this demand to start?	Supports planning and feasibility checks.
44	Is there a required delivery date?	Captures hard deadline or target completion date.
45	What level of delivery are you expecting?	Suggested options: discovery, prototype, proof of concept, MVP, production solution, enhancement, support, or additional capacity.
46	Are there known constraints around timeline, budget, technology, security, procurement, or approvals?	Captures blockers and constraints early.
47	Which teams or stakeholders need to be involved?	Supports stakeholder mapping.
48	Are there dependencies on other projects, vendors, systems, data owners, or decisions?	Captures delivery dependencies.
49	Do you expect NTT DATA to provide advisory support, delivery resources, a managed solution, or a combination?	Supports downstream routing and commercial shaping.

G. Priority and impact

Used to capture customer-stated importance without exposing NTT DATA's internal scoring model.

#	Guided question	Purpose / data captured
50	How important is this demand to your business area?	Captures customer-stated priority.
51	Why is this priority important now?	Prevents unsupported high-priority tagging and captures urgency rationale.
52	What is the expected impact if this is delivered?	Captures impact narrative.
53	What is the impact if this is delayed by three months?	Helps triage urgency and consequence of delay.

54	Are there other related demands that should be considered together?	Supports bundling, grouping, and portfolio-level review.
----	---	--

H. Duplicate and related-demand check

Used after the agent has enough description to search for semantically similar demand items.

#	Guided question	Purpose / data captured
55	Has this idea or a similar request already been raised by your team or another business unit?	Captures user-provided duplicate signal.
56	Would you like me to check for similar demand items before you submit?	Makes duplicate detection transparent, while product rules may still require the check to run before submission.
57	I found a similar demand item. Is your request the same initiative, related but separate, or unrelated?	Captures duplicate handling decision.
58	If related but separate, how is this request different from the existing demand?	Captures differentiation and cross-reference rationale.
59	Should this demand be linked to the similar item for visibility?	Captures customer-safe related-demand linking preference where allowed.

I. Confirmation before submission

Used to ensure the customer explicitly confirms the structured demand record before it is created.

#	Guided question	Purpose / data captured
60	Please review the summary below. Is this an accurate description of your demand?	Human confirmation before creation.
61	Do you want to edit any details before submission?	Allows correction before record creation.
62	Should this be submitted now, or saved as a draft?	Supports incomplete-demand handling.
63	Are you happy for NTT DATA to use this information for triage, scoring, and routing?	Confirms use of submitted information for downstream governance.
64	Do you want to add any final notes, links, examples, or attachments?	Captures any remaining supporting context.

Dynamic question logic

The agent should select the minimum useful set of questions based on the conversation context.

Condition	Agent should prioritise asking
New AI use case	Problem, desired outcome, data sources, AI capability, value hypothesis, benefit KPI, delivery expectation, and duplicate check.
Enhancement to existing solution	Existing solution, current limitation, affected users, expected improvement, urgency, measurable benefit, and dependency on current product owner.
Additional capacity request	Capability or role needed, timeframe, expected start date, reason for capacity need, business impact if not staffed, and related engagement.
Automation opportunity	Process steps, monthly volume, manual time per case, exception rate, rework rate, target time saving, systems involved, and control requirements.
Advisory or discovery request	Business problem, uncertainty areas, decision the discovery should support, expected discovery output, sponsor, target date, and success criteria.
User cannot estimate ROI	Ask proxy questions about volume, time spent, role types involved, error rate, SLA impact, revenue exposure, risk exposure, or external cost.
Duplicate match found	Ask whether the request is the same, related but separate, or unrelated; capture rationale and link where appropriate.
Required fields missing	Ask only for missing mandatory information; if still incomplete, save as draft where supported and explain what remains outstanding.

ROI and benefit fields to populate

The agent should use the answers above to populate or infer structured fields for downstream triage and benefit tracking.

Field	How it is populated
Benefit category	From selected benefits and problem statement.
Estimated annual value	Direct user estimate or calculated proxy using volume, effort, cost, revenue, or risk inputs.
Current baseline	User-provided KPI, SLA, cycle time, cost, error rate, volume, or operational measure.
Target improvement	User-provided desired improvement, percentage reduction, throughput increase, quality uplift, or risk reduction.
Benefit realisation horizon	User-provided expected timing for benefits after delivery.
Confidence level	User-selected confidence: low, medium, or high.
Value evidence	Operational reports, volumes, time estimates, cost data, SLA reports, risk examples, or stakeholder estimate.
ROI narrative	Agent-generated plain-language summary linking the problem, expected outcome, benefit category, value estimate, and confidence level.
Priority input	Strategic importance, urgency, impact of delay, expected value, and feasibility signals. Customer does not see NTT DATA's internal scoring breakdown.

Example ROI follow-up flow — AI contract review request

Customer: We want an AI contract review tool.
Agent: I can help capture that. To understand the value, how many contracts are reviewed each month?
Customer: Around 250.
Agent: Roughly how long does each review take today?
Customer: About 45 minutes for standard contracts, longer for complex ones.
Agent: Which teams are involved in the review?
Customer: Legal Operations and Procurement.
Agent: What improvement would be valuable?
Customer: Reducing review time by 40% and flagging risky clauses earlier.
Agent: Are there costs or risks linked to delayed reviews, missed renewals, external legal review, or compliance exposure?
Customer: Yes, missed renewals and external legal review cost are the main issues.
Agent: Thanks. I will summarise the expected benefits as cycle-time reduction, external-cost reduction, and risk reduction, with a medium-confidence ROI estimate based on monthly contract volume and current review effort.